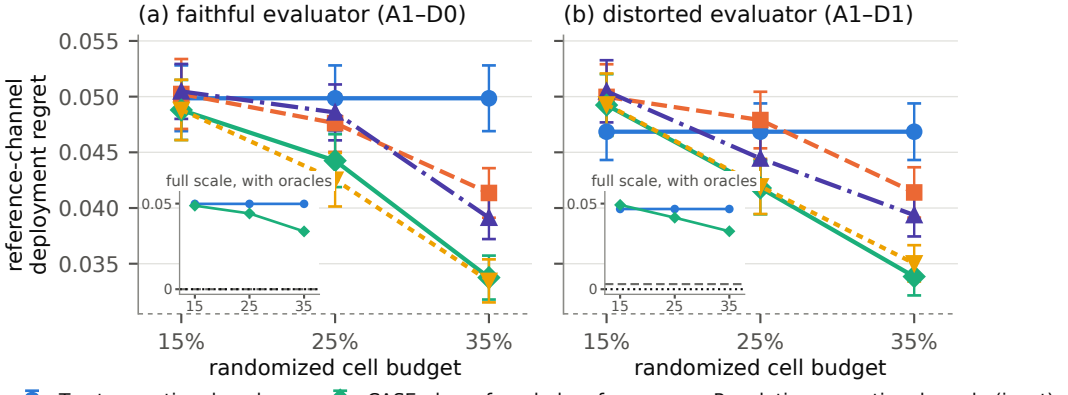
(a) faithful evaluator (A1–D0)
(b) distorted evaluator (A1–D1)
reference-channel deployment regret
0.055
0.050
0.045
0.040
0.035
full scale, with oracles
0.05
0
15
25
35
15%
25%
35%
randomized cell budget
Trust operational eval.
Uniform randomized audit
R-EPIG, observed surface
CASE, deconfounded surface
R-EPIG, deconfounded surface
Population-operational oracle (inset)
Reference-outcome oracle (inset)